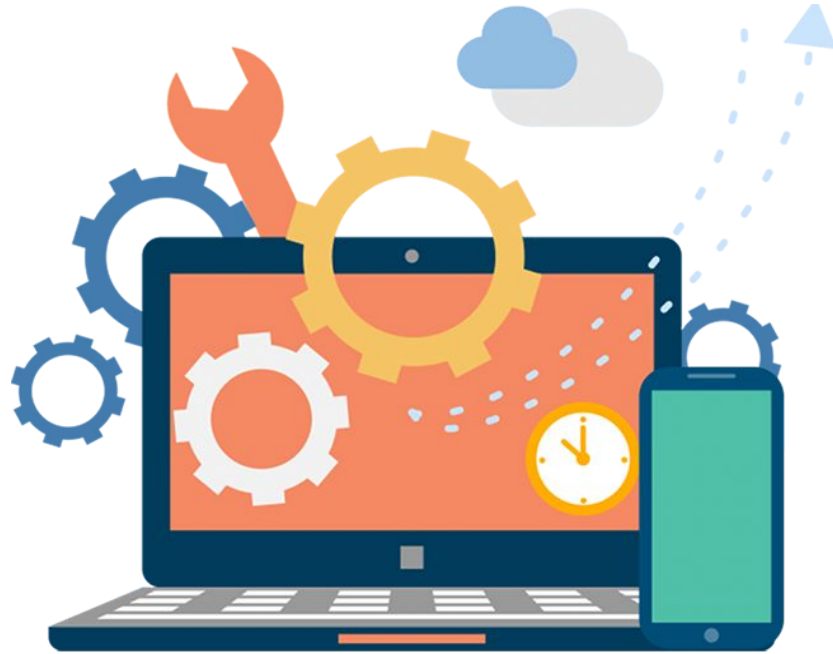


# Python Performance Profiling



# HELLO!

I am **Ikram Ali**

- Data Scientist @ **Arbisoft**
- Working on Deep learning projects for **Kayak**
- [Github.com/akkefa](https://github.com/akkefa)
- [Linkedin.com/in/akkefa](https://www.linkedin.com/in/akkefa)

The logo for Arbisoft, featuring the word "arbisoft" in a lowercase, sans-serif font. The "arbi" part is in blue and the "soft" part is in grey.

# What is Profiling?



# Profiling Definition?

- Measuring the execution time.
- Insight of run time performance of a given piece of code.
- Frequently used to optimize execution time.
- Used to analyze other characteristics such as memory consumption.

# What is Python Profiling?

Measure Performance

# Why Profile?

You can use a profiler to answer questions like these:

- Why is this program slow?
- Why does it slow my computer to a crawl?
- What is actually happening when this code executes?
- Is there anything I can improve?
- How much memory consumed by program?
- How much time taken by each function execution?

# Why You should care about Performance

- “If You Can’t Measure It, You Can’t Manage It.”
- Writing efficient code saves money in modern "cloud economy" (e.g. you need fewer VM instances).
- Even if you don't use clouds, a particular problem domain can have strict performance requirements (e.g. when you have to process a chunk of data in time before the next chunk arrives).

# Available options for measuring Performance in Python



# Command Line

The **time** command is available in \*nix systems.

```
$ time python some_program.py
```

```
real    0m4.536s  
user    0m3.411s  
sys     0m0.979s
```

# Command Line

## PROS

- Easy to use

## CONS

- Very limited information
- Not very deterministic
- Not available on Windows

# Python time Module

Naive approach: **time.time()** statements

```
import time
```

```
initial_time = time.time()
```

```
time.sleep(1)
```

```
final_time = time.time()
```

```
print('Duration: {}'.format(final_time - initial_time))
```

```
Duration: 1.0035021305084229
```

# Python time Module

## PROS

- Easy to use
- Simple to understand

## CONS

- Very limited information
- Not very deterministic
- Manual code modification and analysis

# Python timeit Module

## Better approach: timeit

```
import timeit
```

```
print('Plus:', timeit.timeit("['Hello world: ' + str(n) for n in range(100)]", number=1000))  
print('Format:', timeit.timeit("['Hello world: {0}'.format(n) for n in range(100)]",  
number=1000))  
print('Percent:', timeit.timeit("['Hello world: %s' % n for n in range(100)]", number=1000))
```

Plus: 0.025120729998889146

Format: 0.03536501300004602

Percent: 0.017073806000553304

# timeit Module

## PROS

- Easy to use
- Simple to understand
- Measure execution time of small code snippets

## CONS

- Simple code only
- Not very deterministic
- Have to manually create runnable code snippets
- Manual analysis

# cProfile Module

## Best approach: cProfile

- CPython comes with two profiling tools, profile and cProfile.
- Both share the same API, and should act the same.

```
>>> import cProfile  
>>> cProfile.run('2 + 2')
```

3 function calls in 0.000 seconds

Ordered by: standard name

ncalls	totttime	percall	cumtime	percall	filename:lineno(function)
1	0.000	0.000	0.000	0.000	<string>:1(<module>)
1	0.000	0.000	0.000	0.000	{method 'disable' of '_lsprof.Profiler'}

# You can also run a script with it

```
# slow.py
import time
def main():
    sum = 0
    for i in range(10):
        sum += expensive(i // 2)
    return sum

def expensive(t):
    time.sleep(t)
    return t

if __name__ == '__main__':
    print(main())
```



```
python -m cProfile slow.py
```

25 function calls in 20.030 seconds

Ordered by: standard name

ncalls	tottime	percall	cumtime	percall	filename:lineno(function)
10	0.000	0.000	20.027	2.003	slow.py:11(expensive)
1	0.002	0.002	20.030	20.030	slow.py:2(<module>)
1	0.000	0.000	20.027	20.027	slow.py:5(main)
1	0.000	0.000	0.000	0.000	{method 'disable' of '_lsprof.Profiler'objects}
1	0.000	0.000	0.000	0.000	{print}
1	0.000	0.000	0.000	0.000	{range}
10	20.027	2.003	20.027	2.003	{time.sleep}

# cProfile sort by options

## **ncalls**

For the number of calls

## **tottime**

for the total time spent in the given function

## **percall**

is the quotient of tottime divided by ncalls

## **cumtime**

is the cumulative time spent in this and all subfunctions.

## **percall**

is the quotient of cumtime divided by primitive calls

## **filename:lineno(function)**

provides the respective data of each function

# cProfile result sorted by tottime

```
python -m cProfile -s tottime slow.py
```

25 function calls in 20.015 seconds

Ordered by: **internal time**

ncalls	<b>tottime</b>	percall	cumtime	percall	filename:lineno(function)
10	<b>20.015</b>	2.001	20.015	2.001	{built-in method time.sleep}
1	<b>0.000</b>	0.000	0.000	0.000	{built-in method builtins.print}
1	<b>0.000</b>	0.000	20.015	20.015	slow.py:6(main)
10	<b>0.000</b>	0.000	20.015	2.001	slow.py:13(expensive)
1	<b>0.000</b>	0.000	20.015	20.015	slow.py:3(<module>)
1	<b>0.000</b>	0.000	20.015	20.015	{built-in method builtins.exec}
1	<b>0.000</b>	0.000	0.000	0.000	{method 'disable' of '_lsprof.Profiler' objects}

# cProfile result sorted by ncalls

```
python -m cProfile -s ncalls slow.py
```

25 function calls in 20.015 seconds

Ordered by: **call count**

<b>ncalls</b>	<b>tottime</b>	<b>percall</b>	<b>cumtime</b>	<b>percall</b>	<b>filename:lineno(function)</b>
<b>10</b>	20.020	2.002	20.020	2.002	{built-in method time.sleep}
<b>10</b>	0.000	0.000	20.020	2.002	slow.py:13(expensive)
<b>1</b>	0.000	0.000	20.020	20.020	{built-in method builtins.exec}
<b>1</b>	0.000	0.000	0.000	0.000	{built-in method builtins.print}
<b>1</b>	0.000	0.000	20.020	20.020	slow.py:6(main)
<b>1</b>	0.000	0.000	20.020	20.020	slow.py:3(<module>)
<b>1</b>	0.000	0.000	0.000	0.000	{method 'disable' of '_lsprof.Profiler' objects}

# Easiest way to profile Python code

```
def main():  
    sum = 0  
    for i in range(10):  
        sum += expensive(i // 2)  
    return sum  
  
def expensive(t):  
    time.sleep(t)  
    return t  
  
if __name__ == '__main__':  
    pr = cProfile.Profile()  
    pr.enable()  
    main()  
    pr.disable()  
    pr.print_stats()
```

# cProfile output

25 function calls in 20.030 seconds

Ordered by: standard name

ncalls	tottime	percall	cumtime	percall	filename:lineno(function)
10	0.000	0.000	20.027	2.003	slow.py:11(expensive)
1	0.002	0.002	20.030	20.030	slow.py:2(<module>)
1	0.000	0.000	20.027	20.027	slow.py:5(main)
1	0.000	0.000	0.000	0.000	{method 'disable' of '_lsprof.Profiler'objects}
1	0.000	0.000	0.000	0.000	{print}
1	0.000	0.000	0.000	0.000	{range}
10	20.027	2.003	20.027	2.003	{time.sleep}

# We can also save the output!

```
if __name__ == '__main__':  
    pr = cProfile.Profile()  
    pr.enable()  
    main()  
    pr.disable()  
    pr.dump_stats("profile.output")
```

**How do we use the profiling  
information?**



# pstats Module

- You can use pstats to format the output in various ways.
- pstats provides sorting options. ( **Calls, time, cumulative** )

```
import pstats
```

```
p = pstats.Stats("profile.output")  
p.strip_dirs().sort_stats("calls").print_stats()
```

# pstats module Output

23 function calls in 20.019 seconds

Ordered by: call count

ncalls	tottime	percall	cumtime	percall	filename:lineno(function)
10	20.019	2.002	20.019	2.002	{built-in method time.sleep}
10	0.000	0.000	20.019	2.002	slow.py:14(expensive)
1	0.000	0.000	0.000	0.000	{built-in method builtins.print}
1	0.000	0.000	20.019	20.019	slow.py:7(main)

# An easy way to visualize cProfile results

- Snakeviz library
- PyCallGraph library

# SNAKEVIZ

```
pip install snakeviz
```

```
$ snakeviz profile.output
```

- Snakeviz provides two ways to explore profiler data
- Summaries Times
- You can choose the sorting criterion in the output table

# SNAKEVIZ Browser View

SnakeViz

Reset Root


Reset Zoom

Style: **Icicle** ▾

Depth: **10** ▾

Cutoff: **1** / 1000 ▾

Call Stack



Search:

ncalls	tottime	percall	cumtime	percall	filename:lineno(function)
10	20.01	2.001	20.01	2.001	~:0(<built-in method time.sleep>)
1	0.00011	0.00011	20.01	20.01	slow.py:7(main)
10	7.7e-05	7.7e-06	20.01	2.001	slow.py:14(expensive)
1	6.4e-05	6.4e-05	6.4e-05	6.4e-05	~:0(<built-in method builtins.print>)
1	1e-06	1e-06	1e-06	1e-06	~:0(<method 'disable' of '_Isprof.Profiler' objects>)

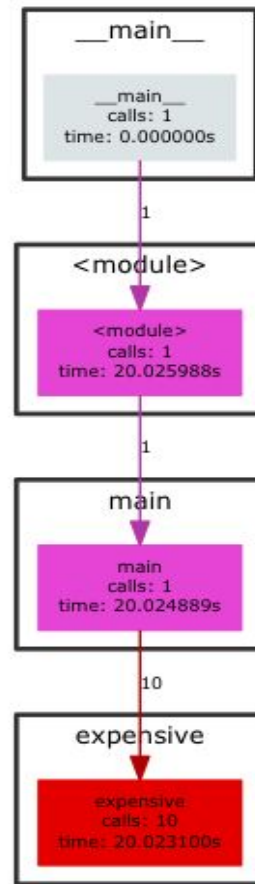
Showing 1 to 5 of 5 entries

# PyCallGraph

pip install pycallgraph

```
$ pycallgraph graphviz -- python slow.py
```

- Visual extension of cProfile.
- Understand code structure and Flow
- Summaries Times
- Darker color represent more time spent



# Other profiling options

## Line profiler

- line\_profiler will profile the time individual lines of code take to execute.
- [https://github.com/rkern/line\\_profiler](https://github.com/rkern/line_profiler)

## Memory profiler

- Monitoring memory consumption of a process.
- line-by-line analysis of memory consumption.
- [https://pypi.org/project/memory\\_profiler/](https://pypi.org/project/memory_profiler/)

# Live Example Interlude

## Profiling Example Code

<https://github.com/akkefa/pycon-python-performance-profiling>





# Thank you.

## Question?

Contact : [mrikram1989@gmail.com](mailto:mrikram1989@gmail.com)